

# Micronavigator: A Comparative Study of Reactive vs. Learning-Based Path Planning

Author: Elsayed Ibrahim

Date: January 3, 2026

Course: Intelligent Robotics

---

## Abstract

This project, titled **Micronavigator**, implements and compares two distinct approaches to autonomous robot path planning in static grid environments: **Potential Fields** (a reactive physics-based method) and **Q-Learning** (a reinforcement learning method). The system was tested across six diverse scenarios, ranging from simple open spaces to complex mazes and large-scale environments. Results demonstrate that while Potential Fields provide computationally efficient,  $O(1)$  reactive planning, they are susceptible to local minima. In contrast, the Q-Learning agent, trained via a Multi-Task Learning architecture, successfully converges on global optimal paths, eliminating local minima at the cost of upfront training time.

---

## 1. Introduction

Path planning is a fundamental problem in mobile robotics, requiring an agent to find a collision-free route from a start point to a goal. This project explores the trade-off between **reactive systems**, which calculate moves on-the-fly based on local gradients, and **deliberative systems**, which learn the global map structure.

The objectives of this project are:

1. To implement a **Gradient Descent Planner** based on Artificial Potential Fields (APF).
2. To implement a **Reinforcement Learning Planner** using Tabular Q-Learning.
3. To develop a robust simulation environment for visualization and benchmarking.
4. To quantitatively compare both methods in terms of path length, execution time, and success rate.

---

## 2. Methodology

### 2.1 Method A: Artificial Potential Fields (Physics-Based)

The robot treats the environment as a continuous landscape of forces. The goal exerts an "attractive" force, while obstacles exert "repulsive" forces. The robot moves by following the negative gradient of the total potential function.

#### Mathematical Formulation:

- Attractive Potential ( $U_{att}$ ): Modeled as a conic well, linearly increasing with distance to the goal.

$$U_{att}(q) = \xi \cdot d(q, q_{goal})$$

- Repulsive Potential ( $U_{rep}$ ): Modeled to rise exponentially as the robot approaches an obstacle within a critical radius  $\rho_0$

$$U_{rep}(q) = \begin{cases} \frac{1}{2}\eta\left(\frac{1}{\rho(q)} - \frac{1}{\rho_0}\right)^2 & \text{if } \rho(q) \leq \rho_0 \\ 0 & \text{if } \rho(q) > \rho_0 \end{cases}$$

- **Total Force:** The robot moves to the neighbor  $q'$  that minimizes

$$U_{total} = U_{att} + U_{rep}$$

#### Local Minima Recovery:

A critical limitation of APF is getting stuck in local minima (e.g., U-shaped traps). The system implements a Oscillation Detection mechanism. If the robot visits the same set of nodes repeatedly, it triggers a Random Walk Recovery mode, where it ignores gradients for 100 steps to "bubble" out of the trap.

### 2.2 Method B: Q-Learning (Reinforcement Learning)

The robot acts as an agent learning to maximize a cumulative reward signal through trial and error. We employ a **Multi-Task Learning** approach, where a single "Universal Brain" is trained on all map scenarios simultaneously.

#### Agent Architecture:

- **State Space (S):** Defined by the tuple (MapID, Row, Col).
- **Action Space (A):** Up, Down, Left, Right.
- **Reward Structure (R):**
  - **Goal:** +1000

- **Collision:** -10
- **Step:** -1 (Encourages shortest path)

Bellman Update Rule:

The agent updates its Q-Table (memory) using the standard Bellman equation:

$$Q(s, a) \leftarrow Q(s, a) + \alpha[R + \gamma \max_{a'} Q(s', a') - Q(s, a)]$$

Implementation Note: The Q-Table uses a dense NumPy array optimized for the Apple M4 Pro architecture, allowing for extremely fast vector lookup compared to standard Python dictionaries.

### 3. Implementation Details

The system is built in Python 3.11 with a modular architecture:

- **core/grid.py:** Handles .txt map parsing and Configuration Space expansion (inflating walls based on robot dimensions).
- **core/navigation.py:** Implements the Gradient Descent logic and local minima detection.
- **ai\_planner.py:** Manages the Q-Learning training loop and inference.
- **main.py:** Provides the real-time GUI for the Potential Field method.

Hardware Optimization:

The project leverages the Apple Silicon (M4 Pro) architecture. By utilizing contiguous memory allocation in NumPy arrays for the Q-Table, the AI training phase achieves high throughput, processing ~15,000 episodes across 6 maps in under 2 minutes.


### 4. Experimental Results

The planners were benchmarked on 6 scenarios, including "Simple", "Maze", "Cluttered", and "Large" environments.

#### 4.1 Comparison: Scenario 3 (Maze)

The "Maze" scenario features a distinct U-shaped trap designed to fail naive planners.

- **Potential Fields:** The robot enters the trap, detects oscillation, enters recovery mode (random walk), and eventually escapes. The path is jagged and suboptimal.
- **Q-Learning:** The trained agent recognizes the trap structure immediately and navigates around it. The path is smooth and optimal.



## 4.2 Quantitative Benchmarks


### Performance Metrics:

- **Success Rate:** Did the robot reach the goal?
- **Path Length:** Total steps taken.
- **Planning Time:** CPU time required to calculate the path.

### Benchmark A: Potential Fields

- **Strengths:** Instant startup (0ms training).
- **Weaknesses:** High variance in path length due to random recovery steps.
- **Result:** Solved all maps, but "Maze" required significantly more steps due to recovery maneuvers.

Micronavigator Benchmark (Robot Size: 2x2)




Benchmark Dashboard - Potential Fields

### Benchmark B: Q-Learning AI

- **Strengths:** Consistent, optimal paths. 100% success rate on trained maps.
- **Weaknesses:** Requires initial training phase.
- **Result:** "Maze" was solved in the theoretical minimum number of steps.

AI Agent (Q-Learning) Performance



Benchmark Dashboard - AI

---

## 5. Discussion

The experiments highlight a fundamental trade-off in robotics:

1. **Reactivity vs. Optimality:** The Potential Field method is excellent for dynamic environments because it recalculates instantly ( $O(1)$ ). However, it lacks "memory," leading to inefficient paths in complex geometries.
  2. **The Local Minima Problem:** While our "Random Walk" recovery successfully solved the Maze, it is non-deterministic. In a time-critical mission, this unpredictability is a risk.
  3. **The Value of Learning:** The Q-Learning agent demonstrated that with sufficient training (15,000 episodes), a robot can "memorize" the topology of the world. The implementation of **Multi-Task Learning** proved that a single model could handle disparate environments (e.g., tight corridors vs. open fields) without retraining.
- 

## 6. Conclusion

**Micronavigator** successfully demonstrated two viable path planning strategies. The project satisfied all requirements, including handling variable robot sizes, visualizing paths, and generating performance statistics.

While Potential Fields remain a robust fallback for unknown environments, the results show that **Reinforcement Learning** provides superior performance in known, static environments, offering a 30-50% reduction in path length for complex scenarios like the Maze.

---

## 7. References

1. Khatib, O. (1986). *Real-time obstacle avoidance for manipulators and mobile robots*. The International Journal of Robotics Research.
2. Sutton, R. S., & Barto, A. G. (2018). *Reinforcement Learning: An Introduction*. MIT press.
3. Scipy Community. *NumPy: The fundamental package for scientific computing*. <https://numpy.org/>